

Municipal Bond Characteristic Factors, Residual Dynamics and Algo Error Correction

Yield-space IPCA K=3 v3 factor model of record, state-space residual signal, and print-level error-correction validation | Research decision note for management, quantitative researchers and traders

Jer-Shen Chen · v3.8 · run muni_ipca_research_v5 (v3.8 notebook) on data 2026-01-05 to 2026-09-30 · October 2026

Abstract

We estimate a three-factor, yield-space Instrumented Principal Components model in which municipal-bond factor loadings are functions of thirteen observable characteristics, and we carry the model through to the quoting engine. On the covered closing-mark panel (1.77 million bond-days, 53,408 CUSIPs, 181 dates), the walk-forward model explains 35.6% of out-of-sample daily yield-change variance with a monthly refit and leaves a point-in-time residual with standard deviation 5.28 bp. Two data rules remove the evaluator artefacts that distorted earlier versions, and the refit-cadence question is closed: monthly, quarterly and frozen Gamma are within 0.3 percentage points of each other. The residual is best described by a pooled state-space filter (near-random-walk drift plus mark noise plus white noise). The filtered drift has daily rank IC +0.083 and a vol-scaled hedged paper Sharpe of 6.8, and the filtered mark noise is a strong mark-quality score (+41.7 bp of trade-versus-mark error per unit, $t = 14.9$). The residual signal predicts where evaluated marks go; it does not predict where trades print relative to the algo quote (-0.14 bp per rank unit, $t = -0.18$).

The deployable result sits one layer down. The algo quote's own pricing error persists from one matched print to the next with coefficient +0.59 ($t = 22.9$), the quote's existing last-value adjustment is already right-sized (-0.037 , so there is no mis-scaled term to fix), and a one-coefficient rule that adds $0.9 \times \text{EWMA}$ of the bond's past errors on the same side of the market (half-life three prints) lowers held-out MAE against the print from 12.95 to 11.60 bp (+1.35 bp, Fama–MacBeth $t = 16.1$, block-bootstrap $t = 11.6$) on 537,845 held-out trades. The side-pooled version of the same rule gains +0.73 bp; keeping the memory on the bond's own side nearly doubles it and turns the short end from a loss into a +2.6 bp gain. A per-side state-space memory, the same rule with the gain and the gap decay learned by maximum likelihood, reaches +0.79 bp, so the fixed three-print half-life is the right object. The gain is concentrated where the factor framework says it should be: single-A callables (+18 to +22 bp in two cells), 4–6 year bonds (+2.0 bp), short bullets (+1.8 bp) and the high slope-beta tercile (+1.5 bp); it is zero or negative in AAA and in the 6–8 year belly. In dollars of price the same-side rule removes about \$2.8 million a month of the quote's \$30 million a month of absolute error on the held-out prints, more than the level model, and \$2.4 million of it in single-A bonds. The same-side rule gains on all three sides (+2.9 bp on dealer purchases, +1.1 bp on dealer sales, +0.4 bp inter-dealer). Five attempts to bring the correction inside the factor framework, a state-space memory on the print clock with and without a per-side state, point-in-time pooling of neighbours' errors in beta space, within-cluster relative value from the residual filter and a per-cluster factor loading, were run under the same protocol and all lose to the same-side memory; the reasons are diagnostic. The algo's error loads on the slope factor (-30.8 bp per unit β_2 , $t = -4.2$) and the correction removes that loading; a direct factor-beta correction does not work on its own (-0.29 bp). A LightGBM level model reaches +1.55 bp but is not yet an object

we would deploy. For the quote band, a fixed width rescaled on a ten-day window covers 78% at 11.2 bp and is more efficient than every conditional band we tried. The recommended stack is therefore the factor model as the common-risk layer, the state-space residual filter as the marks-and-risk layer, and the same-side error memory as the quote layer, gated by the beta-space cells the factor model identifies. The evidence supports deploying that memory behind a shadow period and keeping the level model and conditional bands as challengers.

1. Motivation and Research Question

This study asks three questions in sequence. First, do observable municipal-bond characteristics explain the cross-section of daily evaluated-yield changes well enough to give a stable factor model of record? Second, does the unexplained component carry short-horizon information, and if so about what: the next evaluated mark, or the next trade? Third, where the quoting engine is wrong, does the factor framework tell us where and why, and can a bounded correction be validated against printed trades?

The earlier versions of this work (v1 to v2.3) answered the first question and found the residual's information to be mark-relative rather than quote-relative. The v3 programme accepts that finding and pivots: it keeps the factor model and the residual as the common-risk and marks layer, and tests the quote's own error history as the object to correct. The notebook finds a stable factor model, a residual that is a mark-dynamics object, and a persistent algo error whose correction is statistically and economically meaningful in specific characteristic cells. It does not establish a causal mechanism, and it covers five out-of-sample months with one dispersion cycle.

2. Data

Dataset	Size	Role
Closing-mark panel (Step 3 pipeline output)	1,894,388 rows, 77,379 CUSIPs, 2026-01-05 to 2026-09-30	Target and continuous characteristics
Closing marks, deduplicated	2,680,659 rows, 106,117 CUSIPs	Yield, price, duration lookups
Model panel after filters	1,772,768 rows, 53,408 CUSIPs, 181 dates, L = 13	IPCA fit
Residuals of record (monthly refit)	1,195,848 rows, 44,359 CUSIPs, 123 dates, 7 Gamma versions	All downstream residual tests
AlgoSignal-to-MSRB matched trades	13,741,255 raw rows; 9,645,961 after match filter and trade-id dedup; 6,743,756 in the OOS window	Transaction validation
Trades joined to a prior residual	624,171 trades, 31,493 CUSIPs, 125 trade dates, median residual age 3 days	Signal-to-print tests
Error-correction held-out set	537,845 trades in five monthly folds (May to September)	Baselines A to H and the level model
Band held-out set	344,397 trades in three monthly folds (July to September)	Quote-band calibration

Filters applied before the fit. Rows with a zero duration and a long maturity are excluded (5,792 rows, 0.31%). Yield outliers are flagged per date by a robust z-score (7,048 rows, 0.37%); the target is capped so that 284 rows (0.015%) with a daily move above 250 bp do not drive the moments. Dates with fewer than 1,000 priced bonds are dropped (five dates). Ratings use the final composite rating with a fallback to the composite field; 18.2% of bond-days remain unrated and carry the NR flag.

Coverage boundary. The closing-mark panel covers the OneTick-priced subset of the municipal universe, and all transaction tests are conditional on the trades that match an AlgoSignal quote. Results should be read as statements about the quoted universe, not the whole market. The 20.5% of bond-days whose evaluated mark does not move are retained in the fit and tracked throughout.

2.1 The two data rules

The v5 run showed that two kinds of evaluator days were distorting the model. The v3 notebook makes them explicit rules, fixed before the run.

Rule	Definition	Treatment	Dates caught in this sample
Dispersion day	More than 10% of bonds move more than 10 bp <i>relative to the day's median move</i>	Zero weight in the ALS moments; residuals still produced and flagged	2026-05-27, 2026-09-17, 2026-09-28
Common-move day	Absolute median move above 10 bp	Flagged and kept in the fit	2026-03-20, 03-24, 07-23, 09-10, 09-23, 09-24, 09-28 and others
Partial-mark day	Bonds priced below 60% of the trailing 20-date median count	Dropped	2026-04-22

Why "relative to the median" matters. An evaluator-wide reprice is a day on which many bonds move a lot *in different directions* after the day's common move is taken out: the evaluator is re-levelling individual bonds. A day on which the whole curve moves 15 bp is a market day, and the factor model is exactly the tool meant to explain it. The raw rule used in v4 and v5 (share of bonds with a large absolute move) flagged twelve dates, eight of which were common-move days. Zero-weighting those eight threw away the most informative curve days of September, and that is why the v5 run reported a frozen Gamma beating the weekly refit. With the relative rule the order reverses and the cadence question closes.




Figure 1. The two data rules: share of bonds with a large move relative to the median (top) and the bond count relative to its trailing median (bottom), with the dispersion, common-move and partial-mark dates marked.

3. Research Design and Pipeline

Conceptual orientation only. Quantitative evidence is reported in Section 5.

Step 1 — Data, characteristics and data rules

Construct the bond-date panel from closing marks and lagged security characteristics; apply the dispersion, common-move and partial-mark rules so that the moments the model sees are market days.



Step 2 — Conditional factor fit and refit cadence

Estimate the three-factor yield-space IPCA model by date-weighted alternating least squares. Refit monthly on an expanding window; run quarterly and frozen refits on identical test dates to settle the cadence question.



Step 3 — Residual construction and residual dynamics

Define the point-in-time residual as the observed yield move less the characteristic-conditioned common component. Describe its dynamics with a state-space model (persistent drift, mark noise, white noise) and take the filtered drift as the residual signal of record and the filtered mark noise as a mark-quality score.



Step 4 — Mark tooling

Use the factor path to roll stale marks forward, locate every bond in beta space, and form point-in-time beta-space clusters for the breakdowns that follow.



Step 5 — Transaction validation of the residual

Join the lagged signal to subsequent MSRB prints. Test the signal against the prior evaluated mark and against the existing algo quote with date-aware inference.



Step 6 — Algo error correction and quote band

Build the print-to-print panel of the algo's own error; estimate correction rules on prior months only; validate held-out MAE against the print by age, side, characteristic cell and beta-space cluster; calibrate a quote band around each mid and compare it to a fairly rescaled fixed band.

Research boundary. This workflow maps the thesis and sequencing; it does not itself establish significance, economic value or production readiness. Those rest on the walk-forward design, the residual-age controls, the Fama–MacBeth and block-bootstrap inference, the pre-registered choices in Appendix B, and the cell-level evidence in Section 5.

4. Methodology

4.1 Model

For bond i on date t , the daily closing-yield change in basis points is

$$\Delta y_{i,t} = z_{i,t}' \Gamma f_t + \varepsilon_{i,t} = \beta_{i,t}' f_t + \varepsilon_{i,t}, \quad \beta_{i,t} = \Gamma' z_{i,t} \quad (1)$$

Here $z_{i,t}$ is an $L = 13$ vector of predetermined characteristics, Γ is the $L \times K$ map from characteristics to loadings, f_t is the $K = 3$ vector of latent factor realisations, $\beta_{i,t}$ is the bond-specific loading vector and $\varepsilon_{i,t}$ is the residual. Γ is common across bonds and dates; the cross-section of betas changes as characteristics change. The specification is a conditional decomposition, not a causal model.

4.2 Instruments

Continuous characteristics are lagged one observation and rank-normalised within each date to approximately $[-0.5, 0.5]$, so the inputs are scale-free across dates. OTHER is the omitted geography.

Group	Instruments
Intercept	<i>market_fv</i> = 1
Rate and price structure	Duration ratio (modified duration divided by years to worst), premium/discount, closing-yield level
Term and optionality	Years to worst and extension, defined as years to maturity minus years to worst, floored at zero
Credit	Rating score from the composite rating and an NR flag (18.2% of bond-days)
Liquidity	Trailing 20-observation mark-change rate
Geography	CA, NY, TX and FL indicators (12.2%, 5.4%, 18.8%, 4.5% of rows)

DV01 is excluded, as in v1: its rank correlation with duration is 0.998 and it produced a mechanically offsetting factor with gross leverage near 39. The duration instrument is the ratio to years to worst rather than raw duration, so that it measures convexity and coupon structure rather than repeating the term axis; the largest off-diagonal instrument correlation is 0.49.

4.3 Estimation

Γ is estimated by alternating least squares on per-date cross-products, which avoids materialising the full 1.77-million-row design matrix. Conditional on Γ , each f_t is a cross-sectional least-squares fit with a small ridge term; conditional on the factor path, Γ is updated from the stacked system. Identification follows the IPCA convention: Γ is orthonormalised, factors are ordered by variance and signs are chosen so factor means are non-negative.

Date weighting. Each date's moments enter the Γ update with weight inversely proportional to that date's cross-sectional mean squared target, capped at four times the median weight. Without this, a handful of high-dispersion days contribute most of the squared variation and dominate the estimate of Γ , which is how the June step in the v4 mimicking-portfolio leverage arose: a late-May evaluator day entered the expanding window and re-levelled the factor map. With inverse-variance weights the kept dates span a weight range of 0.08 to 4.0 and the step disappears. Dispersion days receive weight zero.

The number of factors is $K = 3$. In-sample date-weighted variance explained rises from 0.219 at $K = 1$ to 0.244 at $K = 2$, 0.251 at $K = 3$, 0.254 at $K = 4$ and 0.255 at $K = 5$. The third factor is the last one that both adds fit and has an economic name (Section 5.1); the fourth and fifth add 0.3 and 0.2 points of in-sample fit, have no stable loading pattern across Gamma versions and raise the mimicking-portfolio leverage. K sensitivity was closed in v2.3 and is not re-run here.

4.4 Point-in-time discipline and refit cadence

- Γ is refit on an expanding window at month-end, beginning 2026-03-31, giving seven Gamma versions over the out-of-sample window 2026-04-01 to 2026-09-30. Each residual row carries the version that produced it.
- Quarterly and frozen (fit once at 2026-03-31) refits are run on identical test dates so that the cadence comparison is like for like. A dispersion-regime-conditional Gamma was run in v3.0 and dropped once the cadence question closed.
- Continuous instruments are lagged and rank-normalised within the contemporaneous cross-section.
- A residual dated t is computed from the close of day t and is available from the next business day. Transaction tests join each trade to the last residual strictly before the trade's signal time; median residual age is three days.
- Activity buckets are trailing-20 target volatility per bond, lagged one observation and ranked into quintiles within each date, so every bond-day gets a bucket from past data only.
- The error-correction coefficients, the side intercepts, the factor-beta coefficients and the level model are all estimated on months strictly before the test month.
- The beta-space clusters used in the breakdowns are fitted on the first out-of-sample month and applied forward with fixed centroids.

4.5 Factor-mimicking portfolios, the residual-maker and Gamma alignment

Stack the bonds observed on date t into the beta matrix B_t . The factor realisation is a portfolio of that day's yield changes:

$$f_t = (B_t' B_t + \lambda I)^{-1} B_t' \Delta y_t = W_t' \Delta y_t, \quad W_t = B_t (B_t' B_t + \lambda I)^{-1} \quad (2)$$

and the residual is the yield change projected off the beta span,

$$\varepsilon_t = [I - B_t (B_t' B_t + \lambda I)^{-1} B_t'] \Delta y_t. \quad (3)$$

Each column of W_t is the set of bond weights whose return realises one factor, which gives an interpretable factor path, a systematic roll-forward of stale marks and a book exposure vector in model coordinates.

Two details that were wrong before and are fixed here. First, the gross weight of a mimicking portfolio depends on how the factors are rotated: any orthonormal rotation of Γ leaves the fit unchanged but changes each column's gross weight. We therefore report, alongside the per-factor gross weight, the rotation-invariant Frobenius leverage $\sqrt{\text{tr}[(B'B)^{-1}]}$, which is a property of the beta span alone. Second, comparing Γ across refit versions needs the versions to be in the same rotation. Each version is aligned to the first by an orthogonal Procrustes rotation (the orthogonal matrix that brings one Γ closest to another in least squares). The mean across-version standard deviation of the loadings falls from 0.040 raw to 0.020 aligned, which is the stability figure we quote.

4.6 Residual dynamics: the state-space model

The earlier versions modelled the residual with an AR(1). The v4 result that the AR(1) was "dead" (Pearson lag-1 autocorrelation -0.05 while the Spearman was $+0.06$) was not a failure of persistence but a mis-specified model: the residual contains a mark-noise component that reverses one day later and dominates the linear moments, while a small persistent component survives in the ranks. The v3 model of record writes the residual as the sum of three things:

$$r_t = m_t + \eta_t - \eta_{t-1} + \varepsilon_t, \quad m_t = \phi m_{t-1} + \xi_t, \quad (4)$$

where m_t is a slowly moving drift (the persistent rich/cheap state), η_t is mark noise that enters today and leaves tomorrow (an evaluator placing the mark a little off and correcting it), and ε_t is white noise. Four parameters (ϕ and the three standard deviations) are estimated by maximum likelihood on a Kalman filter over 2,000 bonds per fold, with increments winsorised at the 1st and 99th percentiles. The filter then gives, for every bond and date, the best estimate of m_t (the signal of record) and of η_t (the mark-noise estimate used as a mark-quality score).

What the Kalman filter is doing, in plain terms. Each day the filter holds a belief about the bond's hidden drift and how uncertain that belief is. When a new residual arrives, the filter asks how much of it to attribute to a real change in the drift and how much to a transient mark error. If mark noise is large relative to drift changes, a big residual is mostly discounted and the drift estimate barely moves; if the drift is volatile, the estimate follows the residual closely. The weights are not chosen by hand: they fall out of the four estimated variances. The same model, written as a moving-average filter, is the ARMA(1,1) the v4 review asked for, so fitting an ARMA separately would re-estimate the same object with a less interpretable parameterisation; that experiment was run in v2.3 and closed.

Two reference forecasters are kept for comparison: the trailing 20-day mean of the residual (the "level" signal) and a ten-lag ridge filter ("LongConv-lite"). The by-activity-bucket version of the state-space model is run as a diagnostic; the pooled version is the pre-registered signal of record.

4.7 Transaction validation

For each matched MSRB print the mark-relative and quote-relative errors, in basis points, are

$$e^{\text{mark}}_{i,t} = 100 \times (y^{\text{trade}}_{i,t} - y^{\text{prior mark}}_{i,t}), \quad e^{\text{algo}}_{i,t} = 100 \times (y^{\text{trade}}_{i,t} - y^{\text{algo}}_{i,t}). \quad (5)$$

Each is regressed on the signal's within-date percentile rank with controls for log trade size, minutes from signal time and the side-specific MMD spread change. MSRB sides are coded S (dealer sells to customer), P (dealer buys from customer) and D (inter-dealer).

Fama–MacBeth and the block bootstrap. Trades on the same date share market information, so a pooled regression overstates precision. The primary inference runs the regression separately on each of the 125 trade dates and reports the mean slope with a t-statistic from the dispersion of the daily slopes (Fama–MacBeth). Daily slopes can themselves be autocorrelated across adjacent days, so we also resample the daily slopes in five-day blocks 500 times and report the bootstrap t. Where the two disagree, the bootstrap t is the one to believe; in this sample the ratio of bootstrap to Fama–MacBeth t ranges from 0.61 to 1.01, and every headline result survives both.

The mark-quality reading regresses $| \text{trade} - \text{prior mark} |$ on the state-space mark-noise estimate $|\eta|$, and the dealer round trip (P yield minus S yield on the same bond-day) on the same quantity.

4.8 Algo error correction

The quoting engine's yield for a bond is the MMD yield plus a side-specific MMD spread plus a last-value adjustment,

$$y^{\text{algo}}_{i,t} = y^{\text{MMD}}_t + \text{spread}_{\text{side}(i,t)} + \text{adj}^{\text{last value}}_{i,t} \quad (6)$$

and its error at a print is e^{algo} . The print-to-print panel pairs consecutive matched prints of the same bond (120,021 pairs, median gap 3.8 days) and regresses the second error on the first, on the factor move between the two prints and on the residual move between them. The correction rules then use, for each trade, only prints that occurred before its signal time:

- *last_algo_err*: the algo error at the most recent prior matched print, and its age;
- *ewm_algo_err*: an exponentially weighted mean of prior errors with a half-life of three prints;
- *n_prior_prints* and *same_side_as_last*.

The baselines, all estimated on months strictly before the test month:

Label	Mid	What it tests
A	Algo quote	The incumbent
B	$A + \rho(\text{age bucket}) \times \text{last error}$	A one-print memory
C	$A + \rho \times \text{EWMA of past errors}$	A smoothed memory; two parameters
E	A + per-side intercept	Side bias alone
F	$E + \rho \times \text{EWMA}$	The "revised third term": side intercept plus smoothed memory
G	$A + b'\beta$ (factor-beta correction)	Does correcting the factor exposure directly work?
H	$F + b'\beta$	Memory plus factor exposure
D	LightGBM level model with the error features	How much is there to get
D–	The same without the error features	What the error features add

Gains are reported as held-out MAE against the print and as the daily Fama–MacBeth mean of $|e^A| - |e^{\text{mid}}|$, with FM and bootstrap t.

Memories inside the framework (Section 10b of the notebook). The v35 side breakdown showed the pooled EWMA applies the last error from a sale with the wrong sign to the next purchase in the short end. Six further families were run under the same protocol:

Label	Mid	What it tests
S	$A + \rho_s \times \text{EWMA of the bond's prior errors on the same side (pooled EWMA when none)}$	Is the memory side-specific?
S2	$A + \text{side bias} + \rho \times \text{EWMA of errors demeaned by the universe side bias known before the print's month}$	Is a universe side offset enough?
K	$A + \text{side bias} + \rho_K \times \hat{m}$, the Kalman posterior of a local level with daily decay ϕ and diffusion q on the side-adjusted error, observed under noise r , on the bond's irregular print clock, decayed to the trade	The Section 4.6 model applied to the quote's error
P, PS	$A + \rho \times \text{trade-weighted EWMA (half-life 5 days) of the errors of other bonds in the same point-in-time beta-space cluster (PS: same cluster and side), on days strictly before the trade, exact leave-one-bond-out}$	Does the factor model's notion of similarity carry the correction?
CP, KP	Own EWMA plus cluster pooling; state-space memory plus cluster $\times$ side pooling, coefficients jointly	Are the memories complementary?
RV, CRV	$A + b \times (\text{filtered residual drift} - \text{cluster median drift})$; the same on top of C	Does the residual layer's rich/cheap state price the quote's error?
Gc	The direct beta correction estimated per cluster	Is the factor loading of the error local?

The same-side memory is

$$y_{i,t}^S = y_{i,t}^{\text{algo}} + \rho_s \cdot \text{EWMA}_{3 \text{ prints}}(e_{i,\tau}^{\text{algo}} : \tau < t_{\text{signal}}, \text{side}(\tau) = \text{side}(t)), \quad (7)$$

and the state-space memory is the filter

$$m_{i,k} = \phi^{\Delta d} m_{i,k-1} + \xi_{i,k}, \quad \text{Var}(\xi_{i,k}) = q \cdot \Delta d, \quad e_{i,k}^{\text{algo}} - b_{\text{side}} = m_{i,k} + \epsilon_{i,k}, \quad \text{Var}(\epsilon_{i,k}) = r, \quad (8)$$

with k indexing the bond's prints and Δd the days between them; (ϕ, q, r) are maximum-likelihood on prints before the test month and the EWMA is its steady-state special case with $\phi = 1$. The v3.6 version of K carried one state per bond with a universe side offset; the v3.8 version carries one state per bond **and side**, so the side bias is absorbed by the state itself.

Reading guide for the letters. The two tables above are the full glossary. From here on the text narrows to four mids and names them at every mention: **C** the side-pooled EWMA, kept as the reference; **S** the same-side EWMA, the rule of record; **K** the per-side state-space memory; **D** the gradient-boosted level model, the ceiling. **F**, the side intercept plus side-pooled EWMA, was the rule of record in v3.3 to v3.5 and appears only in the historical tables. Every table and figure in Sections 5.4 to 6 carries a one-line legend so no reader has to return here.

4.9 The quote band

For an 80% target, five bands are placed around each mid on the three held-out months: a fixed half-width calibrated on the training months; a fixed half-width rescaled each day on the trailing ten days (the fair benchmark); a raw gradient-boosted quantile band; split conformal, which rescales the quantile band by the empirical 80% quantile of the normalised score on the previous month; and rolling conformal, which does the same on a ten-day window.

Why conformal, and why a rolling fixed band is the fair benchmark. A quantile model's band can be well shaped (wider where errors are larger) and still mis-covered, because the model's quantiles are only as good as its training period. Conformal prediction fixes the level: it computes how large the realised errors were relative to the band on a calibration window and inflates or deflates the band so that the target coverage held there. It guarantees coverage on exchangeable data, which September was not. A conditional band therefore earns its place only if it beats a fixed band that was given the same recalibration opportunity; that is the rolling fixed band. We compare on efficiency, defined as half-width per point of coverage, lower being better.

4.10 Where the gains live

The universe MAE hides the structure the factor model is built to see. Every held-out trade is tagged with its duration bucket, call structure (bullet, callable priced to maturity, priced to call), rating bucket, state, liquidity quintile, factor-beta tercile and point-in-time beta-space cluster (k-means with eight centroids fitted on standardised betas in April and applied forward). Gains are reported per segment and for every duration $\times$ call $\times$ rating cell with at least 2,000 trades. A Fama–MacBeth regression of each mid's error on the three betas asks whether the quote misses a factor axis. The dollar view multiplies the error removed on each trade by $\text{par} \times \text{price}/100 \times \text{modified duration} \times 10^{-4}$, which converts a yield error into dollars of price.

4.11 Pipeline

1. Build the bond-date panel from closing marks and lagged characteristics; apply the three data rules.
2. Estimate date-weighted IPCA with $K = 3$ on an expanding monthly window; keep residuals with Gamma-version lineage; run quarterly and frozen refits on identical dates.
3. Align Gamma versions by Procrustes; build betas, mimicking weights and the rotation-invariant leverage; map the beta space and fit the point-in-time clusters.
4. Fit the pooled state-space model per fold; publish the filtered drift as the signal of record and the filtered mark noise as the mark-quality score; run the vol-scaled hedged paper portfolio as the information metric.
5. Join signals to MSRB prints strictly before the trade; run Fama–MacBeth and block-bootstrap tests against the prior mark and the algo quote.
6. Build the print-to-print algo error panel; estimate baselines A to H and the level model on prior months; report held-out MAE by age, side, segment, cell and cluster, the factor loading of each mid's error, and dollars removed.
7. Calibrate the five bands around the algo quote, the revised third term and the level model; report coverage, half-width and efficiency against the rolling fixed benchmark.

8. Roll stale marks forward with the factor path and report RMSE reduction by horizon and segment.

5. Results

5.1 Factor Structure and Fit

Estimation. The unbalanced panel of 1,772,768 bond-days is estimated by date-weighted alternating least squares with masked observations, warm starts, QR normalisation and a small ridge term. Three dispersion days receive zero weight; the remaining 177 dates carry weights between 0.08 and 4.0.




Figure 2. Market context: the median closing yield with its 10th to 90th percentile band (left) and the cross-sectional standard deviation of the daily yield change (right), with the first out-of-sample date marked. The spikes in dispersion are the evaluator days the data rules address.

Table 1. Factor count and out-of-sample fit

K	In-sample variance explained (date-weighted)	ALS iterations	
1	0.2194	3	
2	0.2439	6	
3	0.2505	7	selected
4	0.2538	7	
5	0.2553	8	

Table 2. Out-of-sample variance explained by month and refit cadence (identical test dates)

Month	Monthly (record)	Quarterly	Frozen at 2026-03-31	Residual sd (bp)	Target sd (bp)
2026-04	0.1512	0.1512	0.1512	7.02	7.62
2026-05	0.1643	0.1600	0.1600	6.48	7.09
2026-06	0.1236	0.1225	0.1225	4.78	5.11
2026-07	0.5022	0.5023	0.5034	3.73	5.29
2026-08	0.4582	0.4583	0.4591	2.64	3.58
2026-09	0.6569	0.6558	0.6480	4.50	7.69
All	0.3561	0.3549	0.3533	5.28	

Table 3. Gamma of record (full sample), loadings sorted by size

Instrument	Factor 1 (level)	Factor 2 (slope)	Factor 3 (yield tilt)	Norm
market_fv	0.939	-0.324	-0.076	0.997
z_closing_yield_lag1	-0.168	-0.226	-0.893	0.936
z_years_to_worst	0.252	0.764	-0.102	0.811
z_extension	0.149	0.477	-0.313	0.590
z_duration_ratio	-0.041	-0.111	0.183	0.218
z_premium_discount_lag1	0.024	-0.132	-0.091	0.162
z_rating_score	-0.009	-0.032	-0.147	0.150
state_CA	-0.015	-0.026	-0.135	0.138
nr_flag	-0.011	-0.009	0.061	0.063
z_liquidity_20, state_NY, state_FL, state_TX				below 0.05
Variance share	50.6%	40.2%	9.3%	
Factor sd (bp/day)	3.37	3.00	1.44	

Paper-reference note. Kelly, Palhares and Pruitt report $K = 3$ out-of-sample individual-bond total R^2 of 45.3% for monthly U.S. corporate-bond excess returns with 30 instruments. The 35.6% here is for *daily* evaluated-yield changes on municipal bonds with 13 instruments, a far noisier target; the two are directionally consistent and not like for like. Our own v1 number (51.2% on a shorter, calmer sample) is not comparable either, because the v3 window includes the April to June period in which the common factor explained little of the daily move.

Metric definitions and interpretation. *Out-of-sample variance explained* is $R^2 = 1 - \Sigma_{(i,t)}(\Delta y_{i,t} - \Delta \hat{y}_{i,t})^2 / \Sigma_{(i,t)} \Delta y_{i,t}^2$ over all out-of-sample bond-days, using the Gamma in force on each date; it is uncentred, so the denominator is squared yield changes around zero. *Residual sd* is the standard deviation of the point-in-time residual in basis points. *Rank IC* is the Spearman correlation, computed within each date, between a signal and the next observed residual of the same bond, averaged across dates; its t-statistic uses the dispersion of the daily values. *Hedged Sharpe* is the annualised Sharpe ratio of a daily long-short paper portfolio in residual space, weights proportional to the signal scaled by each bond's trailing volatility, with the factor exposures hedged out; it is an information metric about the signal, not a tradable return, because evaluated marks are not executable. *MAE gain* is the daily Fama–MacBeth mean of $|\text{algo error}| - |\text{mid error}|$ in basis points, positive when the mid is closer to the print than the algo quote. *Coverage* is the share of held-out prints inside the band and *efficiency* is half-width in bp divided by coverage in points, lower being better.

Finding 1 — Three factors are legible and stable, and the third is the last one worth keeping. Factor 1 is the parallel level move (market_fv 0.94). Factor 2 is the term or slope move: long years-to-worst and high extension bonds move with it, short bonds against it. Factor 3 is a yield tilt: high-yielding bonds move against low-yielding ones, with a secondary callable (extension) and California component. The shares are 50.6%, 40.2% and 9.3%. Rating, liquidity and the state indicators carry loadings below 0.15 and are kept because they are cheap and interpretable, not because they move the fit. Across the seven Gamma versions the aligned loading standard deviation is 0.020. Adding a fourth factor adds 0.3 points of in-sample fit and no stable pattern.




Figure 3. Gamma anatomy: loadings by instrument for the three factors (left) and the factor paths and variance shares (right).

Finding 2 — The refit cadence does not matter, so the cheapest cadence is the right one. Monthly, quarterly and frozen Gamma explain 0.356, 0.355 and 0.353 of out-of-sample variance on identical test dates; no month differs by more than 0.9 points, and September, the only stress month, favours the monthly refit by 0.9 points. The v5 result that a frozen Gamma beat a weekly refit was an artefact of the raw dispersion rule zero-weighting September's curve days (Section 2.1). The pre-registered record is the monthly refit; a weekly refit buys nothing.




Figure 4. Refit cadence: out-of-sample variance explained by month for the monthly, quarterly and frozen Gamma, on identical test dates.

Finding 3 — The fit is regime-dependent and the residual is the quiet-market object. Variance explained is 0.12 to 0.16 in April to June, when the target sd was 5 to 8 bp and most of the daily move was idiosyncratic, and 0.46 to 0.66 in July to September, when the common factors returned. The residual sd falls from 7.0 bp in April to 2.6 bp in August and rises to 4.5 bp in September. The downstream residual tests are therefore tests on a residual whose scale moves by a factor of three; the vol scaling in the paper portfolio and the normalised conformal score exist for this reason.




Figure 5. Fit over time: monthly out-of-sample variance explained, residual and target standard deviation.

Mimicking portfolios. Across dates the gross weight of the three factor portfolios averages 3.26, 2.67 and 2.77 with the aligned Gamma (3.40, 2.26 and 2.97 raw), with net weights 0.60, -0.81 and 0.05 and effective breadth of 4,700 to 6,500 bonds; the factors are realised by broad, modestly levered portfolios, not by a handful of bonds. The maximum $|B'\epsilon|$ across dates is 2.5×10^{-9} , confirming that the residual is orthogonal to the beta span. The June step in the v4 leverage path is gone. The rise in gross weight from August into September appears in the rotation-invariant leverage as well, so it is a real change in the beta span during the stress period, not a rotation artefact.




Figure 6. Gross weight of the mimicking portfolios with the aligned Gamma (left) and with the raw variance-ordered Gamma, where label swaps show as jumps (centre); the rotation-invariant Frobenius leverage (right).

Beta space. The three betas place every bond in a space where distance is factor exposure rather than label. Figure 7 shows the September snapshot; the eight clusters separate cleanly into short bullets, intermediate bullets, 4 to 6 year priced-to-call, 6 to 8 year priced-to-call and long callables, and the rich/cheap ranking inside a cluster produces named comparables (Figure 8). The point-in-time clusters used in Section 5.5 are the April fit applied forward; their labels (duration, share callable, extension) appear in the cluster rows of Table 11.




Figure 7. The beta-space map on 2026-09-30, coloured by duration, call structure, rating and residual.




Figure 8. K-means clusters in beta space (left) and the ten richest and ten cheapest bonds by residual inside cluster C6, the 4 to 5 year priced-to-call group (right).

5.2 Residual Dynamics

Notebook evidence — autocorrelation. On 1.15 million consecutive residual pairs, the lag-1 Pearson autocorrelation is -0.052 and the Spearman is $+0.058$ (Bartlett band ± 0.002); the Spearman stays between $+0.05$ and $+0.08$ through lag 10 while the Pearson is zero from lag 2. Excluding dispersion days changes neither. This is the signature of the model in Section 4.6: a transient mark-noise component that reverses once and dominates the linear moments, and a small persistent component in the body of the distribution.




Figure 9. Residual autocorrelation by lag, Pearson and Spearman, all days and excluding dispersion days.

Table 5. Residual forecasters on five expanding monthly folds (next residual of the same bond)

Forecaster	Rank IC	IC t	Sign accuracy	IC ex dispersion days	Hedged Sharpe (vol-scaled, all)
Trailing-20 mean (level signal)	+0.082	8.0	0.508	+0.080	2.90
LongConv-lite ridge, 10 lags	+0.070	5.4	0.535	+0.072	2.61
State-space, pooled (record)	+0.083	7.5	0.546	+0.086	6.78
State-space, by activity bucket	+0.061	5.7	0.526	+0.065	5.65
Walk-forward selected signal (by prior-month Sharpe)					5.80 over 102 days

Table 6. State-space parameters, pooled, by fold

Fold	ϕ (drift persistence)	sd drift shock	sd mark noise	sd white noise	Log-lik per obs
2026-05	0.943	0.196	0.553	2.338	-2.345
2026-06	1.000	0.002	0.678	2.257	-2.328
2026-07	1.000	0.004	0.542	2.040	-2.209
2026-08	0.989	0.061	0.410	1.949	-2.143
2026-09	0.982	0.067	0.492	1.785	-2.082

Finding 4 — The residual is a near-random-walk drift under mark noise and white noise, and the filter that knows this is the best forecaster. The drift persistence ϕ is 0.94 to 1.00 with a shock sd of 0.00 to 0.20 bp; mark noise has sd 0.4 to 0.7 bp; white noise has sd 1.8 to 2.3 bp. The drift is tiny and almost permanent, which is why the trailing

mean, an unweighted estimate of the same object, is nearly as good a forecaster (IC 0.082 against 0.083) but a much worse portfolio (Sharpe 2.9 against 6.8): the filter's weighting of recent residuals and its handling of unchanged marks give it twice the sign accuracy margin (0.546 against 0.508). The by-bucket filter, pre-registered in v3.0 as the record, is worse (IC 0.061, Sharpe 5.65) and the parameter table says why: in the two most active quintiles the drift persistence collapses to zero and the fit is unstable. The record is the pooled filter; the by-bucket filter stays as a diagnostic.




Figure 10. Linear filters over the residual path: fitted ridge taps against the taps implied by the state-space model (left); monthly out-of-sample rank IC of the pooled filter (centre); state-space parameters by activity bucket, with the drift persistence collapsing in the two most active quintiles (right).

Finding 5 — The body continues and the tails revert. Conditional on the current residual, the next residual continues with a ratio of 0.10 to 0.15 in the central 90% of the distribution and reverses in the top and bottom 1% (continuation -0.05). Excluding dispersion days sharpens both. A linear AR(1) averages over the two regimes and reports nothing; the winsorised filter and rank-based signals see the body. This closes the AR(1) and the raw-ARMA questions from the v4 review.




Figure 11. Two-regime check: mean next residual by current-residual bin, all days and excluding dispersion days.

Finding 6 — The hedged paper portfolio is an information metric, and on that metric the filter is strong. The pooled filter's vol-scaled, factor-hedged portfolio earns 0.32 bp per day of residual with a Sharpe of 6.78 over 121 days

(4.26 unhedged); it is selected in every walk-forward month and the selected-signal Sharpe is 5.80. The Sharpe is 3.9 to 4.6 within every activity quintile, so the result is not carried by stale marks. These are returns on evaluated marks, which cannot be executed, so the figure measures how much the filter knows about where marks go and nothing else.




Figure 12. Paper portfolio: cumulative hedged residual return by forecaster (left) and Sharpe by activity bucket (right).

5.3 Transfer to Transactions

Notebook evidence. 624,171 held-out prints are joined to a residual strictly before their signal time.

Table 7. Fama–MacBeth over 125 trade dates (bp per rank unit)

Score	Target	n	FM slope	FM t	Bootstrap t
One-day residual rank	trade – prior mark	624,171	+2.10	2.9	1.8
One-day residual rank	trade – algo quote	624,171	−0.84	−2.0	−2.0
Trailing-mean rank	trade – prior mark	491,095	−3.13	−4.1	−2.7
Trailing-mean rank	trade – algo quote	491,095	−2.85	−4.6	−3.4
State-space rank (record)	trade – prior mark	410,723	+1.15	1.3	1.2
State-space rank (record)	trade – prior mark, neutralised date × side × size	410,723	+1.27	1.5	
State-space rank (record)	trade – algo quote	410,723	−0.14	−0.2	−0.1
State-space rank, side S	trade – prior mark	138,756	+2.95	2.5	2.4
State-space rank, side P	trade – prior mark	113,274	−1.45	−1.8	−1.7

Table 8. Mark-quality reading

Test	n	FM slope	FM t	Bootstrap t	
trade – prior mark on residual / vol	624,171	+2.59	7.9	5.6	
 trade – prior mark on state-space mark noise η	410,723	+41.7	14.9	11.7	bp per unit
Dealer round trip (P – S yield, same bond-day) on $ \eta $	66,007	+2.24	4.8	3.3	median round trip 5.8 bp

Finding 7 — The residual signal predicts marks, not prints, and the mark-noise estimate is a mark-quality score. Against the prior mark the record signal has a positive slope of +1.15 bp per rank unit that does not reach significance, with the whole effect on the dealer-sale side (+2.95, t 2.5). Against the algo quote it has nothing (–0.14, t –0.2), on all days and excluding dispersion days. The trailing-mean signal is wrong-signed against prints (–3.1 bp, t –4.1), which is the two-regime result of Finding 5 seen through trades: trades print against stale levels. The joint regression of the three scores is collinear and is not reported. By contrast the filtered mark noise $|\eta|$ is a very strong predictor of how far a print will land from the mark (+41.7 bp per unit, bootstrap t 11.7) and of the dealer round trip. The right place for the residual layer is therefore marks and risk tooling (Section 6), not the quote.




Figure 13. Transaction validation: trade minus prior mark by decile of the state-space signal rank (left) and of the trailing-20 mean rank (centre), by MSRB side; Fama–MacBeth slope of trade minus mark on the record signal by residual age (right).

5.4 Algo Error Correction

Notebook evidence — persistence. On 120,021 consecutive matched print pairs (median gap 3.8 days), the algo error at the second print regressed on the first gives +0.586 (FM t 22.9, bootstrap t 17.7). The factor move between the prints enters with –0.08 (t –1.2) and the residual move with –0.22 (t –5.0): prints follow the factor move and ignore the residual move, consistent with Finding 7. Persistence by gap is 0.42 within one day, 0.42 at one to three days, 0.48 at three to seven, 0.37 at seven to fourteen and 0.31 at fourteen to thirty days; it is 0.58 when the two prints are on the same side and 0.51 when they are on opposite sides, so it is not a side bias. 99.9% of held-out trades have a prior matched print before their signal time.

Notebook evidence — the quote's own third term. The algo quote already carries a last-value adjustment. Regressing the error on that adjustment gives –0.037 (t –3.9): a slight over-reaction, –0.06 on the inter-dealer side and 0.00 on dealer sales. The correlation between our last-error feature and the quote's adjustment is +0.09, and +0.06 for the

EWMA. The persistence is therefore information the quote does not use, not a mis-scaled version of a term it already has.

Table 9. Held-out MAE against the print and daily Fama–MacBeth gain over the algo quote (537,845 trades, May to September)

Mid	MAE (bp)	Gain (bp)	FM t	Bootstrap t
A algo quote	12.95			
B algo + $\rho(\text{age}) \times \text{last error}$	13.18	−0.23	−4.0	−3.3
C algo + $\rho \times \text{EWMA error}$	12.26	+0.68	8.1	6.1
E algo + side intercept	12.85	+0.09	4.9	3.3
F = E + $\rho \times \text{EWMA (revised third term)}$	12.29	+0.64	7.0	5.3
G algo + factor-beta correction	13.23	−0.29	−19.0	−13.5
H = F + factor-beta correction	12.23	+0.71	8.5	6.4
D LightGBM level model with error features	11.44	+1.55	21.1	13.7
D− level model without error features	11.76	+1.26	15.8	10.1
Headline excluding duration < 1 year (474,662 trades)	A 8.76 · C 7.77 · F 7.84 · G 9.00 · D 7.55			

Mids: C side-pooled EWMA (reference) · S same-side EWMA (rule of record) · K per-side state-space memory · D gradient-boosted level model · F side intercept + side-pooled EWMA (historical) · A the algo quote.

Table 10. Gain of the pooled EWMA rule C and the same-side rule S by residual age and MSRB side (bp)

Bucket	n	Gain C	t (C)	Gain S
age ≤ 1 day	234,369	+0.90	6.9	+1.64
age 1–3 days	96,105	+0.73	4.4	+1.48
age 3–7 days	104,140	+0.41	2.0	+1.11
age 7–21 days	78,254	+0.64	4.2	+1.19
age > 21 days	24,685	−0.30	−2.3	−0.24
side D (inter-dealer)	202,235	+0.11	1.0	+0.41
side P (dealer buys)	154,995	+1.61	17.9	+2.93
side S (dealer sells)	180,615	+0.55	5.6	+1.10

Mids: C side-pooled EWMA (reference) · S same-side EWMA (rule of record) · K per-side state-space memory · D gradient-boosted level model · F side intercept + side-pooled EWMA (historical) · A the algo quote.

Finding 8 — The algo's error is persistent and a smoothed memory removes a reliable part of it. The coefficient on the side-pooled EWMA estimated on prior months rises from 0.77 in May to 0.84 in September as the training set

grows; the age-bucket coefficients for the single-print rule B (last error $\times \rho$ by age) rise in parallel from 0.3–0.56 to 0.44–0.58. B hurts (−0.23 bp) because one print is noisy; C (side-pooled EWMA) gains +0.68 bp with bootstrap t 6.1 and is positive in every age bucket below 21 days and on both customer sides. The side intercepts add nothing on top: F (side intercept + side-pooled EWMA) is C with noise added, the side biases are P +2 to +4 bp, D −1, S +0.5, and they shrink over the sample. This was the v3.3 result; Finding 11 shows the same memory kept on the bond's own side (S) is the deployable object.

Finding 9 — Correcting the factor exposure directly does not work; correcting the error history does, and it also fixes the factor exposure. G (universe factor-beta correction) is worse than the algo quote (−0.29 bp, t −19) in every segment; added to F (side intercept + side-pooled EWMA) it adds 0.07 bp (H). Yet the algo quote's error loads on the slope factor: Fama–MacBeth of the error on the betas gives +12.5 bp per unit β_1 (t 2.7), −30.8 per unit β_2 (t −4.2) and zero on β_3 . Under F the β_2 loading is +0.8 (t 0.1) and β_1 falls to 8.8; under S (same-side EWMA) it is +4.8 (t 0.8). Under D (level model) it flips to +17.6 (t 3.0), an over-correction. The reading is that the quote mis-prices along the term axis in a way that is bond-specific and persistent; the per-bond error history carries it, a universe-wide beta coefficient cannot.

Finding 10 — The level model shows how much is there, not what to deploy. D (gradient-boosted level model) reaches +1.55 bp in the v3.3 run and +1.57 bp once the same-side memory is one of its features, of which 0.5 bp comes from the error features (D− without them is +1.06). The v5 ablation showed the model depends on the side-specific MMD spread change, and its error loading on β_2 is wrong-signed. It is the ceiling against which the one-coefficient rule is measured, and S (same-side EWMA) gets 86% of the way there with one number.




Figure 14. Error correction, v3.3 run: held-out MAE against the print by residual age for A (algo quote), B (single-print rule), C (side-pooled EWMA), E (side intercept), F (side intercept + EWMA), G, H, D (level model) and D− (left); mean algo error at the current print by decile of the last error, by side (centre); the single-print coefficient by age of the last print (right).




Figure 15. Held-out MAE by month for every mid of the v3.3 run, same letters as Figure 14 (left), and feature importance of D, the level model, with the error features in the list (right).

Notebook evidence — memories inside the framework. All fourteen mids are scored on the same 537,845 held-out trades.

Table 9b. Held-out MAE and gain over the algo quote for the framework memories (bp)

Mid	MAE	MAE ex < 1 y	Gain	FM t	Boot t	P bid	S offer	D	< 1 y
A algo quote	12.95	8.76							
C own EWMA, side-pooled	12.26	7.77	+0.68	8.1	6.1	+1.61	+0.55	+0.11	-1.37
F side intercept + own EWMA	12.29	7.84	+0.64	7.0	5.3	+1.43	+0.46	+0.20	-1.15
S same-side EWMA	11.60	7.52	+1.36	16.3	11.7	+2.93	+1.10	+0.41	+2.62
S2 side bias + side-adjusted EWMA	12.86	8.48	+0.06	0.6	0.4	+0.65	-0.29	-0.07	-1.18
K state-space memory	13.80	9.37	-0.87	-6.7	-4.4	-0.23	-1.64	-0.67	-2.46
P cluster pooling	12.98	8.78	-0.03	-4.1	-3.1	-0.02	-0.06	+0.01	-0.04
PS cluster × side pooling	13.12	9.06	-0.20	-5.4	-2.9	+0.67	-0.89	-0.25	+0.84
CP own EWMA + cluster	12.26	7.77	+0.68	8.1	6.0	+1.61	+0.54	+0.11	-1.39
KP state-space + cluster × side	13.49	9.18	-0.48	-3.1	-1.9	-0.60	-0.33	-0.54	-1.05
RV within-cluster relative value	12.99	8.79	-0.04	-9.2	-4.9	-0.04	-0.06	-0.01	-0.04
CRV own EWMA + relative value	12.27	7.78	+0.67	8.0	6.0	+1.60	+0.54	+0.09	-1.38
Gc local factor loading	15.14	10.93	-2.05	-11.8	-5.8	-1.64	-2.61	-1.80	-2.49
G universe factor loading	13.23	9.00	-0.29	-19.1	-13.5	+0.22	-0.53	-0.44	-0.52
D level model	11.44	7.55	+1.55	21.2	13.7	+2.46	+1.27	+1.14	+3.99

Mids of the v3.6 run: C side-pooled EWMA · F side intercept + side-pooled EWMA · S same-side EWMA · S2 side bias + side-adjusted EWMA · K state-space memory with a universe side offset · P/PS cluster and cluster × side pooling · CP,

KP combinations · RV/CRV within-cluster relative value · Gc per-cluster factor loading · G universe factor loading · D level model.

Table 9c. State-space and pooling parameters by month, estimated on prior months

Test month	ρ (pooled EWMA)	ρ_s (same side)	ϕ per day	implied half-life (days)	q	r	ρ_{PS} (cluster × side)
2026-05	0.77	0.87	0.999	494	1.24	454	0.67
2026-06	0.77	0.87	0.999	826	1.01	459	0.75
2026-07	0.81	0.88	0.999	465	1.11	449	0.79
2026-08	0.82	0.90	1.000	2,193	1.00	468	0.80
2026-09	0.84	0.92	1.000	∞	18.4	326	0.80

Table 9d. The v3.8 run: the four mids kept, on the same 537,845 held-out trades (bp)

Mid	MAE	MAE ex < 1 y	Gain	FM t	Boot t	P bid	S offer	D inter-dealer	< 1 y
A algo quote	12.95	8.76							
C side-pooled EWMA (reference)	12.21	7.74	+0.73	8.7	6.4	+1.67	+0.61	+0.15	−1.04
S same-side EWMA (rule of record)	11.60	7.53	+1.35	16.1	11.6	+2.93	+1.10	+0.39	+2.60
K per-side state-space memory	12.17	7.99	+0.79	8.8	5.6	+2.53	+0.41	−0.17	+1.17
D level model, with the same-side memory as a feature	11.42	7.55	+1.57	20.1	12.9	+2.44	+1.35	+1.14	+4.19
D− level model without the error features	11.96	7.99	+1.06	12.6	8.0	+1.94	+0.93	+0.52	+2.79
Per-side state-space parameters by test month	ϕ per day 0.997, 0.996, 0.998, 0.998, 0.994	half-life 212, 172, 406, 383, 108 days	q 1.0, 1.0, 1.5, 1.0, 10.4	r 285, 305, 333, 337, 239	ρ_K 1.52, 1.47, 1.34, 1.38, 1.20				

Mids: C side-pooled EWMA (reference) · S same-side EWMA (rule of record) · K per-side state-space memory · D gradient-boosted level model · F side intercept + side-pooled EWMA (historical) · A the algo quote.

Finding 11 — The memory is side-specific, and keeping it on the bond's own side doubles the gain and repairs the short end. S (same-side EWMA) gains +1.36 bp against +0.68 for C (side-pooled EWMA) in the v3.6 run, and +1.35 against +0.73 in the v3.8 rerun of Table 9d, with bootstrap t 11.7, and the coefficient rises from 0.87 to 0.92 as the

training window grows. It is positive on every side (+2.9 bp on the bid, +1.1 on the offer, +0.4 inter-dealer), in every age bucket below 21 days, and in the sub-one-year bucket (+2.6 bp, where the pooled rule lost 1.4 bp). The universe side intercept (S2) does not reproduce this: the round trip that flips the sign is cell-specific and the only object that carries it is the bond's own error on that side. The headline MAE excluding sub-one-year bonds falls from 8.76 to 7.52 bp; the level model, with its full feature set, reaches 7.55 on the same rows. One coefficient on the bond's own same-side history captures what the gradient-boosted model captures.

Finding 12 — The factor-integrated memories all lose to the same-side memory, for reasons that are diagnostic.

The v3.6 state-space memory K, one state per bond with a universe side offset, is worse than the quote (−0.87 bp). Its parameters say why: ϕ is estimated at 1.000 and r at 450 bp² against q near 1, so the filter is a very slow random-walk level under heavy observation noise, which is the pooled EWMA with a longer window and the same side-mixing defect, and the shrinkage coefficient of 1.1 to 1.4 says it under-reacts. The v3.8 version with one state per bond and side (Table 9d) repairs the sign on the bid and in the short end and reaches +0.79 bp, above C (side-pooled EWMA, +0.73) but well below S (same-side EWMA, +1.35): the learned half-life of 100 to 400 days is far longer than the three-print memory the data reward, and the shrinkage coefficient of 1.2 to 1.5 again says the filter under-reacts. The fixed short memory is the right object and K is retired. Cluster pooling adds nothing (P −0.03): the neighbours' trailing error is contemporaneously correlated with the bond's error (the cluster × side coefficient is a stable 0.67 to 0.80) but it does not predict it a day later, because the cell-level component of the error moves with the factors between prints; only the bid and the sub-one-year bucket gain from it (+0.67, +0.84). Within-cluster relative value from the residual filter is wrong-signed ($b = -2.5$ to -6.0 , gain −0.04), consistent with Finding 7: the residual knows about marks, not prints. The factor loading estimated per cluster is the worst mid (−2.05 bp): the slope-factor loading of the error is real at the universe level but too noisy to estimate locally. The two layers therefore divide the work. The factor model and its residual filter own common risk, marks and the map of where the quote is wrong; the quote's own same-side error history owns the correction.




Figure 16. Memories inside the framework, v3.6 run: gain over the algo quote with a 95% band for every mid in Table 9b (left) and the gain by MSRB side (right).




Figure 17. Gain by duration bucket and MSRB side, v3.6 run, for F (side intercept + side-pooled EWMA), S2 (side bias + side-adjusted EWMA), K (state-space memory with a universe side offset), PS (cluster $\times$ side pooling) and KP: only a side-specific memory repairs the short end.




Figure 17b. Gain by duration bucket and MSRB side, v3.8 run, for C (side-pooled EWMA), S (same-side EWMA), K (per-side state-space memory) and D (level model). S is positive on the bid and the offer in every bucket from 1 to 11 years; the inter-dealer short end is the one cell it still loses.

5.5 Where the Gains Live

The universe gain of +0.68 bp hides a very uneven map, and the unevenness is the point: the factor framework identifies the cells where the quote is wrong in a persistent way.

Table 11. Gain over the algo quote by segment (bp, FM t), S same-side EWMA and D level model, v3.8 run

Segment	Share of trades	Algo MAE	Gain S (same-side EWMA)	t	Gain D (level model)
Rating A	14%	18.8	+5.05	11.0	+4.19
Rating AA	72%	12.6	+0.95	13.0	+1.38
Rating AAA	12%	8.7	−0.06	−0.7	+0.15
Bullets	37%	8.6	+1.75	12.8	+1.67
Callable, priced to maturity	9%	11.5	+1.11	2.6	+0.79
Priced to call	55%	16.1	+1.13	11.3	+1.66
Duration 1–2.5 y	24%	11.4	+1.19	7.2	+1.33
Duration 2.5–4 y	15%	8.0	+1.13	8.7	+1.38
Duration 4–6 y	20%	9.3	+2.42	14.4	+2.28
Duration 6–8 y	24%	5.8	+0.00	0.0	+0.01
Duration 8–11 y	4%	12.6	+3.56	3.9	+3.09
Duration < 1 y	12%	44.4	+2.60	6.4	+4.19
Liquidity Q2	20%	14.2	+2.22	10.7	+2.21
State CA	17%	14.6	+1.63	10.1	+2.05
β_3 tercile high	33%	10.6	+1.86	14.0	+1.95
β_3 tercile low	33%	14.1	+1.00	6.9	+1.47
β_1 tercile low	33%	22.6	+1.93	10.5	+2.56
Cluster C0 (1.1 y bullets)	13%	13.9	+2.78	2.8	+1.85
Cluster C4 (1.5 y bullets)	12%	9.1	+1.60	9.3	+1.77
Cluster C1 (7 y callables)	33%	7.5	+0.93	6.3	+0.89
Cluster C7 (1.4 y callables, extension 5.9)	6%	45.7	+1.34	3.5	+2.43
Side P, dealer buys (bid)	29%	12.7	+2.93	26.7	+2.44
Side S, dealer sells (offer)	34%	13.0	+1.10	11.0	+1.35
Side D, inter-dealer	38%	13.1	+0.39	3.3	+1.14

Mids: C side-pooled EWMA (reference) · S same-side EWMA (rule of record) · K per-side state-space memory · D gradient-boosted level model · F side intercept + side-pooled EWMA (historical) · A the algo quote.

Table 12. Top and bottom duration × call × rating cells by the gain of S, same-side EWMA (cells with at least 2,000 trades), v3.8 run

Cell	n	Algo MAE	MAE S	Gain S (same-side EWMA)	t	Gain D (level model)
8–11 y, callable to maturity, A	3,129	39.2	15.6	+23.6	5.9	+18.5
4–6 y, priced to call, A	7,357	32.6	12.7	+19.9	11.7	+17.5
1–2.5 y, bullet, A	11,169	17.1	12.2	+4.9	3.3	+2.6
2.5–4 y, priced to call, A	3,958	12.1	9.0	+3.1	3.5	+3.2
2.5–4 y, bullet, AA	35,009	8.2	6.1	+2.1	8.3	+2.4
1–2.5 y, bullet, AA	60,085	10.6	8.5	+2.1	12.3	+2.4
< 1 y, priced to call, AA	46,527	45.0	43.1	+1.9	6.8	+4.2
4–6 y, priced to call, AA	47,766	8.8	7.2	+1.6	8.0	+1.6
1–2.5 y, priced to call, AAA	4,323	10.4	11.1	−0.7	−5.7	−0.6
1–2.5 y, priced to call, AA	28,978	11.6	12.6	−0.9	−9.0	−0.7
4–6 y, callable to maturity, AA	2,533	10.6	11.7	−1.0	−5.0	−1.1
8–11 y, callable to maturity, AA	8,417	8.4	10.5	−2.1	−2.2	−1.0

Mids: C side-pooled EWMA (reference) · S same-side EWMA (rule of record) · K per-side state-space memory · D gradient-boosted level model · F side intercept + side-pooled EWMA (historical) · A the algo quote.

Finding 13 — The gains sit in single-A callables, intermediate bonds and short bullets, and the quote is already right in AAA and the 6–8 year belly. Under S (same-side EWMA) the two A-rated callable cells see the quote's MAE fall from 33 to 39 bp to 13 to 16 bp, a gain of 20 to 24 bp on 2% of trades. Single-A bonds as a group gain 5.1 bp (t 11.0), 4 to 6 year bonds 2.4 bp (t 14.4), bullets 1.8 bp (t 12.8) and the high yield-tilt tercile 1.9 bp (t 14.0). In AAA (−0.06 bp) and the 6 to 8 year belly (0.00 bp) the quote is already well behaved and the correction is neutral; the two 8 to 11 year callable-to-maturity AA cells and the 1 to 2.5 year priced-to-call cells are the places it hurts. The sub-one-year and near-call segments have quote MAEs of 40 to 120 bp, yield-space noise on bonds whose yield is barely defined; the same-side memory gains 2.6 bp there where the side-pooled C lost 1.0 bp. They dominate the universe MAE and should be reported separately, which is why the headline excluding duration below one year (8.76 to 7.53 bp) is the number to quote.

Finding 14 — The residual signal's information is also concentrated, in the short end. Within clusters the record signal's rank IC is 0.10 to 0.15 in the two sub-one-year clusters and 0.07 in the short callables, against 0.01 to 0.05 elsewhere, and the within-cluster hedged Sharpe is highest in the short callables C7 (4.9) and the 4 year callables C3 (3.7). Against prints the within-cluster slopes are small and mixed. The two layers divide the map: the residual layer knows most about the short idiosyncratic end, where the quote's own error is noise; the error-correction layer works in the intermediate and single-A cells, where the quote's error is persistent.




Figure 18. Where the gains live, v3.8 run: gain of S (same-side EWMA) by duration $\times$ call structure (left); gain by beta-space cluster for S and D, the level model (centre); Fama–MacBeth loading of each mid's pricing error on the three factor betas, A algo quote, S, D (right).




Figure 19. The achievement map, v3.8 run: the ten best segments across every axis (left) and the ten best and five worst duration $\times$ call $\times$ rating cells with the MAE before and after (right). Bars are S, the same-side EWMA, with a 95% band from the Fama–MacBeth t ; the diamond is D, the level model.



Figure 20. Gain of S (same-side EWMA, bars) and D (level model, diamond) by characteristic, one panel per axis: duration, call structure, rating, MSR side, state, liquidity and beta-space cluster. v3.8 run.




Figure 21. Gains in factor coordinates for S (same-side EWMA, green) and D (level model, grey): level, slope and yield-tilt beta terciles, point-in-time within each trade date, with the algo MAE and the share of trades under each tick. v3.8 run.

Notebook evidence — by MSR side. The same gains split by the side of the print, P a dealer purchase from a customer (the bid), S a dealer sale (the offer), D inter-dealer.

Table 12a. Gain of S, same-side EWMA, by characteristic and side (bp; * marks $|FM\ t| \geq 2$; D, the level model, in brackets), v3.8 run

Segment	P bid	S offer	D inter-dealer
All trades	+2.93* (D +2.44)	+1.10* (D +1.35)	+0.39* (D +1.14)
Rating A	+6.67* (D +5.05)	+4.59* (D +3.67)	+4.33* (D +4.05)
Rating AA	+2.70* (D +2.39)	+0.71* (D +1.18)	-0.22* (D +0.79)
Rating AAA	+0.85* (D +0.35)	-0.24 (D +0.06)	-0.58* (D +0.09)
Bullets	+2.34* (D +1.72)	+1.61* (D +1.50)	+1.31* (D +1.83)
Callable to maturity	+2.00* (D +0.52)	+1.14* (D +1.05)	+0.71 (D +0.75)
Priced to call	+3.49* (D +3.24)	+0.77* (D +1.34)	-0.13 (D +0.88)
Duration 4–6 y	+2.89* (D +2.33)	+1.90* (D +2.10)	+2.53* (D +2.45)
Duration 8–11 y	+3.98* (D +2.55)	+3.69* (D +3.23)	+3.30* (D +3.26)
Duration 6–8 y	+0.25* (D -0.04)	-0.15* (D +0.04)	-0.12 (D -0.06)
Duration < 1 y	+12.37* (D +12.3)	+2.81* (D +2.80)	-2.66* (D +0.97)
Cluster C5 (0.3 y bullets)	+18.67* (D +19.5)	+0.70 (D +0.80)	-4.25* (D -0.13)
Cluster C7 (1.4 y callables)	+9.98* (D +8.84)	+0.44 (D +1.02)	-2.62* (D +0.04)
Liquidity Q2	+3.78* (D +3.17)	+1.86* (D +1.84)	+1.37* (D +1.84)

Mids: C side-pooled EWMA (reference) · S same-side EWMA (rule of record) · K per-side state-space memory · D gradient-boosted level model · F side intercept + side-pooled EWMA (historical) · A the algo quote.

Finding 15 — With the same-side memory the correction is positive on the bid and the offer almost everywhere; the inter-dealer short end is the one cell it still loses. Under F (side intercept + side-pooled EWMA, v3.5) the sub-one-year bucket and the two short clusters gained 4 to 7 bp on the bid and lost 3 to 8 bp on the offer and inter-dealer, because a pooled memory applies a sale's error with the wrong sign to the next purchase. Under S (same-side EWMA, Table 12a) the offer side is repaired: +2.8 bp in the sub-one-year bucket, +0.7 and +0.4 in the two short clusters, and the bid gains 12 to 19 bp there. Single-A, bullets and the 4 to 11 year callables gain on all three sides. The residual losses are inter-dealer prints in the short end (-2.7 bp below one year, -4.3 and -2.6 in the short clusters), AAA inter-dealer (-0.6) and the belly on the offer (-0.15): an inter-dealer print has no customer side, so a same-side memory has less to say about it. The shadow gating is therefore by cell and by side, with inter-dealer excluded below one year.




Figure 22. Gain of S (same-side EWMA) by characteristic and MSRB side, v3.8 run: duration, call structure, rating and beta-space cluster against P bid, S offer and D inter-dealer, with an asterisk at |FM t| ≥ 2.




Figure 23. The five best duration × call × rating cells on each side under S (same-side EWMA, bars) with D (level model, diamond), with the MAE before and after. v3.8 run.




Figure 24. The record residual signal within beta-space clusters: within-cluster hedged Sharpe (left) and next-day rank IC (right).

5.6 The Quote Band

Table 13. Bands around the algo quote, 80% target, three held-out months (344,397 trades)

Band	Coverage	Mean half-width (bp)	Efficiency (bp per coverage point)
Fixed (calibrated on training months)	75.0%	10.4	0.139
Rolling fixed (rescaled on trailing 10 days)	78.2%	11.2	0.143
Raw quantile GBM	69.9%	14.8	0.212
Split conformal (previous month)	70.1%	14.8	0.211
Rolling conformal (trailing 10 days)	77.3%	18.1	0.234
Around the level model D, rolling fixed	78.1%	11.4	0.146
Around the level model D, rolling conformal	77.9%	15.9	0.204
Around S (same-side EWMA), fixed	76.5%	11.5	0.150
Around S (same-side EWMA), rolling fixed	78.5%	11.8	0.150
Around S (same-side EWMA), rolling conformal	77.9%	15.6	0.200

Mids: *C* side-pooled EWMA (reference) · *S* same-side EWMA (rule of record) · *K* per-side state-space memory · *D* gradient-boosted level model · *F* side intercept + side-pooled EWMA (historical) · *A* the algo quote.

Finding 16 — A rescaled fixed band beats every conditional band we tried, around every mid. The fixed band collapses in September (63.8% coverage) because the training months were calmer. Rolling conformal recovers coverage to 77% and is flat across width deciles (0.73 to 0.83), which is what conformal is for, but its mean half-width is 16 to 18 bp and its top width decile averages 70 bp. The fair benchmark, a fixed width rescaled on the same ten-day window, covers 78% at 11.2 bp with no model at all. On efficiency the rolling fixed band is 0.143 to 0.150 against 0.200 to 0.235 for rolling conformal around every mid, including S (same-side EWMA), where the conditional band is at its best (0.200) and still loses. The conditional band's shape is right and its level is too wide; until the quantile model is sharper, current practice (a fixed width, now rescaled on a short window) is the right band.




Figure 25. The quote band around the level model: coverage by predicted-width decile for the five bands (left); daily coverage under monthly and rolling calibration, with the rolling conformal scale (centre); mean half-width by mid and side, rolling fixed against rolling conformal (right).

5.7 Roll-Forward of Stale Marks

Table 14. Rolling a stale mark forward with the factor path (RMSE, bp)

Horizon (days)	n	Stale RMSE	Rolled RMSE	RMSE reduction
1	1,195,848	6.6	5.3	20%
3	1,108,043	12.1	8.4	31%
5	1,023,100	16.5	10.6	36%
10	838,230	23.8	14.2	40%
At 5 days, by segment				
Duration 1–2.5 y	131,923	13.4	5.9	56%
Duration 2.5–4 y	128,270	12.7	4.1	68%
Duration 4–6 y	162,238	14.0	4.0	71%
Duration 6–8 y	133,688	14.6	4.2	71%
Duration 8–11 y	42,443	13.7	5.0	64%
Duration < 1 y	401,175	20.2	15.8	22%
Callable to maturity	98,757	15.2	7.0	54%
Priced to call	380,942	14.9	8.8	41%
Bullets	543,401	17.8	12.1	32%

Finding 17 — The factor model moves intermediate and long marks; the short end is idiosyncratic. For bonds with one to eleven years of duration, rolling a five-day-stale mark forward with the factor path removes 56 to 71% of the RMSE. Below one year, which is 40% of bond-days, it removes 22% and drags the universe figure to 36%. This is the strongest single argument for the factor model as marks tooling, and it is already in the identified cells.




Figure 26. Rolling a stale mark forward with the factor path: MAE by horizon, stale against rolled (in-panel upper bound).




Figure 27. Roll-forward at five days by duration bucket, call structure and beta-space cluster.

5.8 External Benchmarks and Interpretation

Corporate-bond IPCA reference. Kelly, Palhares and Pruitt's *Modeling Corporate Bond Returns* applies the same characteristic-conditioned latent-factor architecture to monthly U.S. corporate-bond excess returns and finds, as we do, that a small number of characteristic-driven factors explains most of the common variation and that the loadings are stable. Our daily yield-change target is noisier and our 13-instrument set is smaller; the magnitudes are not comparable, the structure is.

Attention-factor and muni relative-value references. The attention-based statistical-arbitrage literature motivates learning economically relevant similarity across assets rather than imposing fixed peer groups; our beta-space clusters and within-cluster rich/cheap ranking follow that principle with the IPCA betas as the learned coordinates. The muni relative-value literature motivates the residual as a conditional fair-value state; the v2.3 peer relative-value signal built on beta-space neighbours was tested and closed because it added nothing to the state-space filter. None of these papers validates the transfer to MSRB prints or the error-correction result; those rest on the tests in Sections 5.3 to 5.5.

6. Application to the Market-Making Model

Audited result. Three outputs of the programme have cleared the research gates and one has not.

1. **The factor model of record** ($K = 3$, 13 instruments, monthly refit, two data rules) is stable across versions and cadences and explains the common move in the months when there is one. It is ready as the common-risk layer.
2. **The residual layer** is a marks-and-risk object: the pooled state-space filter predicts the next mark (IC 0.083, hedged Sharpe 6.8), the mark-noise estimate predicts how far prints land from the mark (bootstrap t 11.7), and the factor roll-forward removes 56 to 71% of stale-mark RMSE for 1 to 11 year bonds. It does not predict prints relative to the quote and should not be used as a quote charge.
3. **The algo error correction S** (same-side EWMA: quote + $0.9 \times$ EWMA of the bond's past errors on the same side, half-life three prints) lowers held-out MAE by 1.35 bp universe-wide with bootstrap t 11.6, is positive on all three sides and in the short end, gains 20 to 24 bp in the two single-A callable cells, 5 bp in single-A overall and 2.4 bp in 4 to 6 year bonds, and removes the quote's slope-factor loading. It removes \$2.8 million a month of pricing error on the matched prints. It is ready for a bounded shadow deployment in those cells. The factor-integrated memories (state-space with and without a per-side state, cluster pooling, relative value, local loading) were tested under the same protocol and lose to it.
4. **The conditional band** is well shaped and too wide; the rescaled fixed band is current best practice. The level model (+1.55 bp) remains a challenger.

Decision. Deploy S, the same-side EWMA, as a shadow third term in the cells of Table 12 with positive, significant gain (rating A; duration 1 to 6 and 8 to 11 years; bullets; the high yield-tilt tercile), with the correction capped and logged. Table 12a sets the side gate: all three sides in single-A, bullets and the 4 to 11 year callables; bid and offer but not inter-dealer below one year and in the two short clusters; nothing in the 6 to 8 year belly, AAA inter-dealer, the 1 to 2.5 year priced-to-call cells and the 8 to 11 year callable-to-maturity AA cell. The recommended stack is the factor model (common risk), the state-space residual filter (marks and risk) and the same-side memory (quote), with the factor model's beta-space cells as the gate. Promote the state-space outputs (filtered drift, mark-noise score, factor roll-forward, beta-space comparables) as research and desk tooling with no pricing impact. Keep the level model and the conditional band as challengers behind the same gates. Do not use the residual signal as a quote adjustment.

6.1 Market-Making Evaluation Gates

Gate	Question	Required evidence	Status for S, the same-side EWMA
1. Signal validity	Is the error memory stable?	Persistence by gap and side pair; coefficient path over months; FM and block-bootstrap t	Passed: persistence 0.31 to 0.58 in every gap bucket and both side pairs; ρ_s rises monotonically from 0.87 to 0.92; bootstrap t 11.6
2. Incremental fair value	Does it beat the quote on the print?	Held-out MAE by age, side, segment and cell against A; the quote's own last-value term checked; factor loading of the error	Passed: +1.35 bp universe-wide, positive on all sides and in every age bucket below 21 days; neutral in AAA and the 6 to 8 year belly; the quote's own term is right-sized (-0.037); the cell and side map is Tables 12 and 12a
3. Execution utility	Does bounded use improve quoting?	RFQ replay with fills, adverse selection, inventory, turnover and tail loss, in the shadow cells	Not yet run: the shadow deployment is the test
4. Economic size	What is it worth?	Par- and duration-weighted dollars of error removed per month, by cell	Measured: \$2.8 million a month, 9% of the quote's absolute error, \$2.4 million of it in single-A bonds (Section 6.3)

6.2 Controlled Integration Path

- Shadow third term.** Compute S (same-side EWMA) alongside the live quote for every trade; log the correction, the same-side EWMA state, the age and the cell; no pricing impact.
- Cell gating.** Apply the correction only in cells where the trailing-quarter gain is positive and significant, re-evaluated monthly on the same held-out protocol; start with the cells in Table 12.
- Bounded challenger.** Cap the correction at a multiple of the cell's recent error sd; exclude sub-one-year and near-call bonds; carry the rolling fixed band as the width.
- Production gate.** Require sustained net utility in the RFQ replay, stable calibration of the band, and no deterioration of the factor loading of the corrected error before the term enters the quoting model.

Candidate mapping for testing, not deployment:

$$y_{i,t}^{\text{quote}} = y_{i,t}^{\text{MMD}} + \text{spread}_{\text{side}} + \text{adj}_{i,t}^{\text{last value}} + 1[c \in C^*] \cdot \text{clip}[\rho_s \cdot \text{EWMA}_3(e_{i,\tau}^{\text{algo}} | \text{same side}), -\text{cap}(c), +\text{cap}(c)] \quad (9)$$

where C^* is the set of validated characteristic cells, ρ_s is estimated on prior months (0.92 at the end of this sample) and $\text{cap}(c)$ is set per cell. The EWMA is over prior matched prints of the same bond on the same side of the market, so the term is zero for a bond with no same-side history and the rule degrades to the incumbent quote (or, by choice, to the side-pooled memory).

6.3 Economic Size

The error removed on each held-out trade is converted into dollars of price with $\text{par} \times \text{price}/100 \times \text{modified duration} \times 10^{-4}$. The par filter drops 1,392 of 537,845 trades with non-finite, non-positive or implausible par; the remaining trades carry \$45.5 billion of par over five months.

Table 15. Dollar error removed on held-out prints (five months, 536,453 trades, \$45.5 billion par), v3.8 run

Mid	\$ removed per month (\$k)	Par-weighted gain (bp)
Algo quote absolute error, for scale	30,254 per month	
C side-pooled EWMA (reference)	+2,565	+0.56
S same-side EWMA (rule of record)	+2,844	+0.48
K per-side state-space memory	+836	−0.20
D level model	+2,571	+0.92
By segment, S		
Rating A (14% of trades, \$6.3 bn par)	+2,402	
Rating AA (72%, \$32.5 bn)	+572	
Rating AAA	−139	
Duration 4–6 y	+1,923	
Duration 8–11 y	+748	
Duration 6–8 y	+17	
Priced to call	+1,870	
Bullets	+541	
Top cells, S		
4–6 y, priced to call, A (7,356 trades, \$469 mm)	+1,505	+28.5
8–11 y, callable to maturity, A (3,125 trades, \$173 mm)	+997	+28.9
4–6 y, priced to call, AA (47,715 trades, \$3.1 bn)	+432	+1.4
6–8 y, priced to call, AA (72,940 trades, \$8.0 bn)	+365	+0.3
2.5–4 y, bullet, AA	+286	+1.8

Mids: C side-pooled EWMA (reference) · S same-side EWMA (rule of record) · K per-side state-space memory · D gradient-booster level model · F side intercept + side-pooled EWMA (historical) · A the algo quote.

Audited result. S (same-side EWMA) removes about \$2.8 million a month of the quote's \$30 million a month of absolute pricing error on the prints that matched, 9% of the total, more than D (level model, \$2.6 million) and C (side-pooled EWMA, \$2.6 million); K (per-side state-space memory) removes \$0.8 million. The par-weighted gain of S (0.48 bp) is below C's (0.56) because S's extra gains sit in short, low-duration bonds that carry little dollar duration, while its dollar total is higher because it also gains in the 4 to 11 year callables. The dollars are concentrated: single-A bonds

account for \$2.4 million of the \$2.8 million, and the two single-A callable cells alone remove \$2.5 million a month on \$640 million of par, a par-weighted gain of 29 bp. The 6 to 8 year belly, which carries the most par (\$15 billion) and the most absolute error (\$59 million), is flat. By month the rule is slightly negative in May, when the coefficient was estimated on one month of data, and reaches \$1.3 million in August. This is an accuracy measure, not a P&L: how much of it a desk captures depends on which quotes are hit, which is the fill model the RFQ replay in Gate 3 supplies.




Figure 28. Dollar error removed per month by mid, C side-pooled EWMA, S same-side EWMA, K per-side state-space memory, D level model (left), by beta-space cluster (centre), and the ten cells that remove the most dollars under S (right). v3.8 run.

7. Limitations and Next Steps

- The results are predictive associations on five out-of-sample months with one dispersion cycle; September is the only stress month. The error-correction coefficient rises every month as training data accumulate, so the universe gain is probably understated, but one quarter of out-of-sample is one quarter.
- The factor model's out-of-sample fit is regime-dependent (0.12 to 0.66 by month). Downstream residual results are results about a residual whose scale moves by a factor of three; the vol scaling and the normalised band score are the controls.
- The residual signal's slope against the prior mark is positive and not significant at the universe level; its information is in the short idiosyncratic end and on the dealer-sale side. It is a marks object and is presented as one.
- Rating is the latest composite value with a fallback rather than a point-in-time history; its loading is small.
- The universe MAE is dominated by sub-one-year and near-call bonds with 40 to 120 bp errors. All headline numbers should be quoted excluding duration below one year, and the cell view is the honest view.
- The level model depends on the side-specific MMD spread change and over-corrects the slope loading; it is a ceiling, not a candidate.
- The conformal band's coverage guarantee assumes exchangeability, which September violated; the rolling version restores coverage at the cost of width.
- Every gain in this note is an absolute distance to the print. It does not distinguish a mid that moves toward the print from one that crosses through it by a smaller amount, and the two are not worth the same to a dealer. A signed decomposition (toward, crossed by less, crossed by more), a crossing rate by side and an asymmetric loss are the first additions for the next version.

- The inter-dealer short end is the one segment where the same-side memory loses; inter-dealer prints have no customer side, so the gate excludes them below one year.
- The dollar view is an accuracy measure on matched prints, not a P&L.

Next quarter. Add the direction-aware view of the gain (signed decomposition, crossing rate by side, asymmetric loss); run the shadow same-side memory in the identified cells and collect the RFQ replay evidence for Gate 3; rerun the notebook monthly so the coefficient path and the cell map are refreshed; sharpen the quantile model (print features at signal time, cell indicators) so that the conditional band can compete with the rescaled fixed band on efficiency; add a point-in-time rating history; and extend the out-of-sample window through the next dispersion cycle.

7.1 Bringing the correction back inside the factor framework: what was tried and what is next

The v3.6 notebook ran the three extensions proposed after v35 under the Section 10 protocol. The results (Table 9b) settle two of them and reshape the third.

1. **State-space error memory.** Specified as a local level on the side-adjusted error with a universe side offset (v3.6), it loses to the quote (-0.87 bp): $\phi = 1.000$, $r \approx 450$, $q \approx 1$ make it a slow random-walk level under heavy observation noise with the side-mixing defect of the pooled EWMA. Specified per bond and side (v3.8, K in Table 9d) it reaches $+0.79$ bp, above C (side-pooled EWMA) and well below S (same-side EWMA). Its steady state is S, so it could only add by learning the gain and the gap decay, and it learns a half-life of 100 to 400 days where three prints is what the data reward. Closed.
2. **Cluster-level pooling.** Neighbours' trailing errors are strongly correlated with a bond's error on the same day (coefficient 0.67 to 0.80, stable) but do not predict it across days (P -0.03 bp), because the cell-level component moves with the factors between prints. Pooling helps only where the own history is thin or noisy: the bid ($+0.67$) and the sub-one-year bucket ($+0.84$). It remains the only correction available for a bond with no print history and is worth keeping as the fallback in that case.
3. **Within-cluster relative value.** Wrong-signed and closed ($b = -2.5$ to -6.0 ; gain -0.04). The residual layer's rich/cheap state is about marks, not prints, exactly as Finding 7 found at the universe level.
4. **Local factor loading.** The worst mid (-2.05 bp). The slope-factor loading of the quote's error is a universe-level fact that a per-bond memory removes; estimating it per cluster adds noise.

Appendix A. Closed Experiments

Each line is a hypothesis that was tested in an earlier version with a pre-stated metric and closed with a number. They are not re-run in v3.

Experiment	Version	Result	Why closed
Linear AR(1) residual forecast	v4	Rank IC -0.085 raw; wrong-signed	Mark noise dominates linear moments; replaced by the state-space model
ARMA(1,1) via AR(∞) taps	v2.3	Same object as the state-space filter, worse parameterisation	Redundant
LongConv ridge filter	v4, v2.3	Rediscovered the trailing mean; IC 0.070	Kept only as a reference forecaster
Mark shrinkage toward the factor fit	v4	No improvement in trade-versus-mark error	Closed
Residual-driven quote width (wider/thinner)	v4	Under-covered (71%) and no MAE gain	The residual does not predict print dispersion; replaced by conformal and rolling fixed bands
Quote skew sizing on the trailing-mean signal	v5	-0.10 bp MAE	Economically negligible
Peer relative value in beta-space neighbourhoods	v2.3	Adds nothing to the filter	Closed
Factor hedging of flow exposure	v2.3	No measurable effect	Closed
K sensitivity (K = 2, 4)	v2.3	K = 4 adds 0.3 points in-sample, unstable loadings	K = 3 fixed
Weekly refit	v3.0	Within 0.3 points of monthly	Monthly fixed
Dispersion-regime-conditional Gamma	v3.0	No gain once the relative dispersion rule was in place	Dropped
By-bucket state-space filter as record	v3.0	IC 0.061 against 0.083 pooled; ϕ collapses in active quintiles	Pooled filter fixed as record
Single-print error correction B	v3.1	-0.23 bp	One print is too noisy; EWMA kept
Side intercepts E	v3.1	$+0.09$ bp, shrinking	Not worth a parameter
Direct factor-beta correction G	v3.3	-0.29 bp (t -19)	The mis-pricing is bond-specific; the error history carries it
Joint regression of the three residual scores	v3.1	Collinear (coefficients $+6.5$, -11.1 , $+7.2$)	Not reported
Side bias + side-adjusted EWMA (S2)	v3.6	$+0.06$ bp	A universe side offset does not carry the cell-specific round trip; the same-side memory does
State-space error memory with universe side offset (K, KP)	v3.6	-0.87 , -0.48 bp; $\phi = 1.000$, $r/q \approx 450$	Slow random-walk level with side mixing
Per-side state-space error memory (K, v3.8)	v3.8	$+0.79$ bp against $+1.35$ for the same-side EWMA; half-life 100 to	Its steady state is the same-side EWMA; learning the decay only adds

Experiment	Version	Result	Why closed
		400 days, ρ_K 1.2 to 1.5	noise
Cluster pooling of neighbours' errors (P, PS, CP)	v3.6	-0.03, -0.20, +0.68 (= C)	Contemporaneous, not predictive across days; kept as the no-history fallback
Within-cluster relative value from the residual filter (RV, CRV)	v3.6	-0.04 bp, $b = -2.5$ to -6.0	Residual knows marks, not prints
Per-cluster factor loading (Gc)	v3.6	-2.05 bp	Universe-level fact, too noisy locally

Appendix B. Pre-Registered Choices

Fixed on 2026-10-07 (v3.1, after the v31 review) and not changed by the v32 or v33 results.

Choice	Value
Number of factors	3
Refit cadence of record	monthly
Level-signal window; activity window	20; 20
Record residual signal	pooled state-space filter (the by-bucket filter is a diagnostic)
EWMA half-life	3 prints
Revised third term	algo + per-side intercept + $\rho \times \text{EWMA}(\text{past algo errors})$, all estimated on prior months
Conformal rolling window; band benchmark	10 days; fixed width rescaled on the same window
Dollar metric	$(\text{algo error} - \text{mid error}) \times \text{par} \times \text{price}/100 \times \text{modified duration} \times 10^{-4}$, per month
Breakdown axes	duration, call structure, rating, state, liquidity, factor-beta terciles, 8 beta-space clusters fitted on the first OOS month
Dispersion rule	share of $ \text{target} - \text{median}_t > 10 \text{ bp}$ above 10% of bonds $\rightarrow$ zero weight
Common-move rule	$ \text{median}_t > 10 \text{ bp} \rightarrow$ flagged and kept
Partial-day rule	bonds priced < 60% of trailing-20-date median $\rightarrow$ dropped
Date weighting	inverse variance, cap $4\times$
Duration instrument	ratio to years to worst
Portfolio scaling	volatility
Band coverage target	80%
v3.6 cluster memory	trade-weighted EWMA of other bonds' algo errors in the same PIT beta-space cluster, half-life 5 days, days strictly before the trade
v3.6 state-space memory	local level with daily decay ϕ and diffusion q on the side-adjusted algo error per print, observation noise r ; ML on prints before the test month

Appendix C. Version Lineage

Version	What changed	What it settled
v1, v1.1	Yield-space IPCA K = 3; transaction audit; quantile-charge validation	Residual information is mark-relative, not quote-relative; no charge model clears the gates
v4	Procrustes alignment; two-horizon signal; MSRB sides fixed; signal-time cut	AR(1) is dead; level model gain falls to 1.1 bp after the signal-time cut; 20% unchanged marks; the June leverage step
v5	Date weighting; state-space model; print-to-print panel; ablations; conditional band	Step gone; SSM best forecaster; algo error persistence 0.585; level model depends on the MMD spread change; band under-covers
v2.3	ARMA, peer RV, flow hedging, K sensitivity, sizing, bootstrap, robustness	All closed with numbers (Appendix A)
v3.0	Clean rebuild; two data rules; cadence study; by-bucket filter as record; error-correction baselines	Data rules catch 04-22 and 05-27; cadence closed
v3.1	Pooled filter as record; last-value diagnostic; baselines E, F; rolling conformal; PIT clusters; segment breakdown	Quote's own term right-sized; EWMA rule +0.68 bp; the gains have a map
v3.2	Beta-tercile and cluster breakdowns; roll-forward by segment; factor loading of the error	Error loads on the slope factor; F removes it
v3.3	Baselines G, H; rolling fixed band benchmark; dollar view with finite par	Direct beta correction fails; rescaled fixed band wins on efficiency
v3.4	Section 12a achievement map: top segments, per-axis grid, factor terciles, scorecard	The "where it works" exhibits come straight from the notebook
v3.5	MSRB side as a breakdown axis; side × characteristic; top cells per side; dollar view on real data	Bid-side result outside bullets and single-A; short end flips sign by side; \$2.5 million a month removed, 90% single-A callables
v3.6	Section 10b: same-side and side-adjusted EWMA's, state-space memory, cluster pooling, within-cluster relative value, per-cluster loading	Same-side memory +1.36 bp, positive on all sides and in the short end; every factor-integrated memory loses to it
v3.7	The same-side memory carried into the breakdowns, the side heatmaps and the dollar view	Superseded by v3.8 before a run
v3.8	Closed mids and the by-bucket and LongConv forecasters removed; per-side state-space memory; same-side memory as the mid of record in Sections 11 and 12; stage cache keyed on config, data fingerprints and code; vectorised print-history memories	S +1.35 bp, \$2.8 million a month, positive on bid and offer almost everywhere; per-side K +0.79, retired; a cold run computes six cached stages in about six and a half minutes and a warm run reuses them

Artifacts: `artifacts_v3/` (residuals of record, Gamma versions raw and aligned, algo error-correction panel, conformal band panel, beta-space snapshot, results registry). Notebook: `muni_ipca_research_v3.ipynb` on branch

claude/nifty-wright-j93nps . *Reviews of each run:* REVIEW_V4_RESULTS.md , REVIEW_V5_RESULTS.md ,
REVIEW_V31_RESULTS.md , REVIEW_V32_RESULTS.md .